

# Learning Energy Efficient Jumping Strategies for Flexible-Legged Systems

Andrew Albright \* Joshua Vaughan \*\*

\* University of Louisiana at Lafayette, Lafayette, LA 70503 USA  
(andrew.albright1@louisiana.edu)

\*\* University of Louisiana at Lafayette, Lafayette, LA 70503 USA  
(joshua.vaughan@louisiana.edu)

**Abstract:** Legged locomotive systems have many advantages over their wheeled counterparts, such as their ability to navigate rough terrain. They have the ability to deploy many techniques to overcome obstacles, one of which is jumping. Still, there are disadvantages to overcome when using legged systems, such as their lack of energy efficiency. To combat this lack of efficiency, flexible links can be used to conserve energy that would otherwise be wasted during locomotion. Furthermore, control methods can be utilized which improve both a jumping system's ability to jump high and their ability to conserve power. Using reinforcement learning (RL) to create these controllers for flexible-legged jumping systems can lead to control strategies that maximize jump height while minimizing power usage.

**Keywords:** Power Efficient, Reinforcement Learning, Flexible Robots, Legged Locomotion, Jumping

## 1. INTRODUCTION

Legged locomotive robots have many advantages over their wheeled counterparts, for example their ability to more easily navigate harshly uneven terrain (Park et al. (2017); Blackman et al. (2018); Seok et al. (2015)). However, there are disadvantages as well, one of which is power consumption. There are several factors that contribute to this disadvantage. For example, in many cases several motors are required to actuate several linkages to move only a few feet, whereas a wheeled system needs only a single motor to turn a wheel to accomplish the same task. Another factor, one that is particularly prevalent when navigating uneven terrain, is the challenge of defining how a walking robot places its feet in such a way that reduces wasted footfall energy, *while maintaining stability*.

In an effort to alleviate the power consumption issues seen when using legged systems, research has been conducted which replaces rigid aspects of said systems with flexible ones. It has been shown that this not only leads to higher performance but also higher efficiency (Sugiyama and Hirai (2004); Galloway et al. (2011); Hurst (2008); Seok et al. (2015)).

While the introduction of flexible components solves some challenges related to performance measures like power consumption, it raises other challenges. Modeling the systems becomes more difficult because the models become highly nonlinear. Different approaches have been taken to mitigate these issues, a popular and successful example being the use of series elastic actuators instead of flexible links (Pratt and Williamson (1995); Iida et al. (2005); Ahmadi and Buehler (1997)). Other solutions seek to create control methods which are suited to non-rigid systems. One of those control methods is reinforcement learning

(RL), which creates controllers based on learned control strategies that are developed through interactions with an environment. These RL controllers are typically called control agents or just agents.

In this work, RL is used to develop control strategies for a simplified jumping robot. The RL agent is trained to maximize jump height while conserving power. ~~It is shown that when tasked as such, the agent learns more precise control strategies to maximize the jumping performance of the system, while also balancing power usage.~~

The next section will look at some related work in the fields of RL, flexible systems, and power conservative-control strategies. In Section 3, a description of the environment used for training and evaluating the RL agents will be provided. A short description of the RL algorithm used for this work is presented in Section 4. Then, in Section 5, a breakdown of the reward functions used to train the agents is provided. The experiments performed will be discussed in Section 6, and evaluations are provided in Section 7. Then finally, in Section 8, conclusive remarks will be made.

## 2. PREVIOUS WORK

### 2.1 Reinforcement Learning for Legged Locomotion

Research has shown that using RL for defining control strategies for legged systems is a viable path for simple and even complex legged locomotive systems. In the case of using a model-based method, reductions of the number of environment interactions during controller training can be achieved by an order of magnitude when compared to model-free methods (Yang et al. (2019)). However, model-based methods expose other challenges. Learning a model often requires many interactions, and the inevitably

This is your  
a running  
at the power

Do you have  
for this?...  
but every  
other choice

limited environment that is learned leads to limited exploration during controller training. As such, model-free methods are often of interest as they can be deployed directly on hardware to learn based on interactions with the environment or trained initially in simulation on a non-perfect model, then moved to hardware to finish training. Specifically, the viability of using RL to train different agents to locomote while navigating uneven terrain has shown great success (Peng et al. (2016)). Similar successes have been seen using a host of different environments, including (Reda et al. (2020)). In addition, they were also able to display the apparent difficulties of defining an environment for creating robust controllers.

## 2.2 Reinforcement Learning for Flexible Systems

As discussed, flexible systems do have some advantages over their rigid counterparts, particularly ones which are used for locomotion tasks. Typically, for model-based RL methods used for flexible systems, learning a controller is necessary as the model of the system is usually difficult to create (Thuruthel et al. (2019)). In support of mitigating the negatives that are associated with learning a model of an environment, model-free methods have been deployed for flexible systems and show the potential of controllers that do not require a model of the system (Dwiel et al. (2019)). Additionally, comparing RL control strategies to more traditional control strategies such as PD control, it has been shown that RL is practically applicable for controlling flexible systems (He et al. (2020)). Furthermore, RL control can be used for partial control of flexible systems. An example being the use of an RL controller to damp out vibration while a separate controller is used to determine general pose (Cui et al. (2019)).

## 2.3 Flexible Robots Improved Performance

Using flexible components within robotic systems has shown great potential for conserving power. Of the different approaches taken, a popular and proven technique is the use of flexible joints (series-elastic actuators) to increase energy efficiency (Pratt and Williamson (1995); Ahmadi and Buehler (1997)). Flexible joints are not the only way to increase efficiency, a similar technique which seeks to emulate the tendons seen in nature can produce similar results (Folkertsma et al. (2012)). In the work completed by Seok et al. (2015) and Seok et al. (2013), design principles for efficient quadruped locomotion are discussed, and evidence supports that the use of flexible components increases efficiency.

## 2.4 Control for Power Efficiency

Rather than relying solely on the mechanical design of a system to increase efficiency, control techniques can be employed to increase power conservation. An example being the use of model predictive control methods that lead to more efficient locomotion strategies (Harper et al. (2019)). In contrast to studying traditional methods of control to accomplish energy efficient strategies, little work has been done to evaluate the potential of using RL control strategies to directly effect power efficiency.

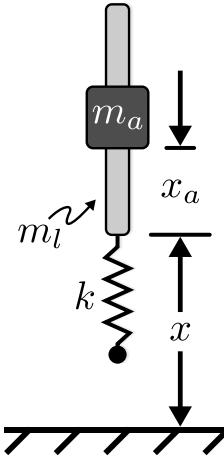


Fig. 1. Pogo-stick Model for a Simplified Flexible-Legged Jumping System

Table 1. Training and Evaluation Parameters

Model Parameter	Value
Mass of Leg, $m_l$	0.175 kg
Mass of Actuator, $m_a$	1.003 kg
Natural Frequency, $\omega_n$	11.13 Hz
Spring Constant, $k$	200000 N/m
Actuator Stroke, $(x_a)_{\max}$	25 mm
Actuator Velocity, $(\dot{x}_a)_{\max}$	2.0 m/s
Actuator Acceleration, $(\ddot{x}_a)_{\max}$	10 m/s

## 3. POGO-STICK ENVIRONMENT

### 3.1 The Pogo-stick Environment

A simplified model of a flexible-legged jumping system is the pogo-stick model shown in Fig. 1. This model can be used as a base representation for many different jumping animals, which in turn represent many different kinds of jumping gaits (Blickhan and Full (1993)). As such, this model is used as the environment that RL agents are tasked with learning control strategies for. Launching the pogo-stick into the air requires the agent to accelerate the mass,  $m_a$ , along the rod,  $m_l$ .

The masses of the actuator and leg are represented by  $m_a$  and  $m_l$ , respectively. A nonlinear spring with constant  $k$  is used as a flexible component. A damper (not shown in Fig. 1) is ~~noted~~ parallel to the spring and is represented by the variable  $c$ . The vertical position relative to the ground is represented by  $x$ , and the position which the actuator moves along the rod is represented by  $x_a$ . The values of many of these parameters are detailed in Table 1. Note that the system is constrained to move vertically so that the agent does not learn to keep the system balanced.

The equations of motion for the system are:

$$\ddot{x} = \alpha \left( \frac{k}{m_t} x^3 + \frac{c}{m_t} \dot{x} \right) - \frac{m_a}{m_t} \ddot{x}_a - g \quad (1)$$

where  $x$  and  $\dot{x}$  are position and velocity of the rod respectively,  $\ddot{x}_a$  is the acceleration of the actuator, and  $m_t$  is the mass of the complete system. Ground contact

determines the value of  $\alpha$ , so that the spring and damper do not supply force while the leg is airborne:

$$\alpha = \begin{cases} -1, & x \leq 0 \\ 0, & \text{otherwise} \end{cases} \quad (2)$$

#### 4. REINFORCEMENT LEARNING

The algorithm used to define the weights  $\phi$  is Stable Baselines implementation of Twin Delayed Deep Deterministic Policy Gradient (TD3) (Fujimoto et al. (2018)). This is an actor-critic, off-policy RL algorithm wherein the actor is represented by the policy  $\pi_\phi$  taking the actions  $a$ , and the critic is represented by the estimated expected return of taking action  $a$  in state  $s$  and following the policy  $\pi$  from then after. The critic is a neural network itself with parameters  $\theta$  that are updated according to temporal difference error found between a set of twin target networks. These target networks are updated to follow the critic network every defined  $n$  updates of the critic network.

#### 5. REWARD FUNCTIONS

##### 5.1 Rewarding the Agent for Jumping High

To accomplish the task of maximizing jump height, a reward function which represents the height of the pogo-stick above the ground at any given time step was used. The signal returned to the agent was normalized with a predefined maximum height so that if the pogo-stick was at or above that maximum, the agent received maximum reward. This reward function is defined as:

$$r = \frac{x_t - x_{min}}{x_{max} - x_{min}} \quad (3)$$

where  $x_t$  is the height at the current time step,  $x_{min}$  is the minimum height of the system, and  $x_{max}$  is the max height of the system, which is set based on expectations gathered from experimental data. In this work, it is set to 0.9 meters.

##### 5.2 Rewarding the Agent for Jumping Efficiently

The second reward function used is one that seeks to balance jumping height with power consumption. To do this, the reward signal returned at each time step is the ratio of the current height to the power used from time zero to the current time step. The power used at a given time step is defined by:

$$p_t = m_a \ddot{x}_{a_t} \dot{x}_{a_t} \quad (4)$$

where  $m_a$  is the mass of the actuator,  $\ddot{x}_{a_t}$  is the acceleration of the actuator at time  $t$ , and  $\dot{x}_{a_t}$  is the actuator velocity at time  $t$ . The efficiency at a given time step is then:

$$e_t = \frac{x_t}{\sum_{t=0}^t p_t} \quad (5)$$

This efficiency value is normalized where the maximum and minimum limits are determined based on experimental

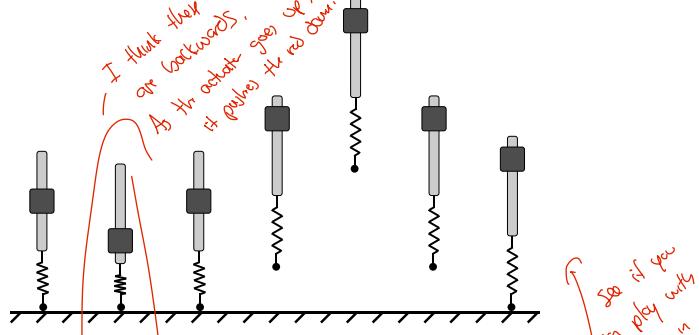


Fig. 2. Pogo-stick Single Jump

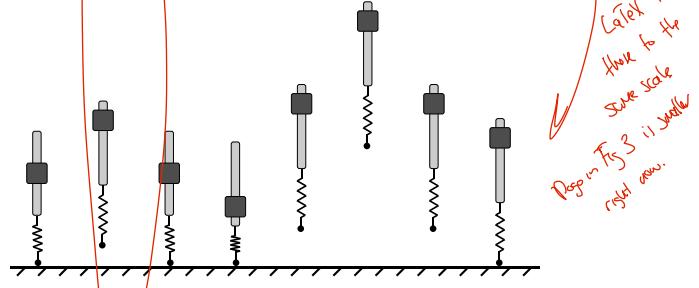


Fig. 3. Pogo-stick Stutter Jump

realizations of efficiency expectations. For this work, the maximum and minimum efficiency values are set to  $e_{max} = 0.002$  and  $e_{min} = 0$  respectively. The reward function for the case of maximizing agent efficiency can then be written as:

$$r = \frac{e_t - e_{min}}{e_{max} - e_{min}} \quad (6)$$

To receive any reward the agent is required to first leave the ground. Maximizing reward requires the agent to maximize jump height while minimizing power used.

#### 6. EXPERIMENTS

##### 6.1 Jumping Types

To evaluate if reinforcement learning, and more specifically if the TD3 algorithm (Fujimoto et al. (2018)) can define more efficient control strategies, two types of jumping experiments are performed. For the first, the agent is tasked with learning both high jumping and efficient jumping strategies for single jumps with the pogo-stick environment, like the one shown in Fig. 2.

For the second jumping strategy, the agent is tasked with learning both high jumping strategies and efficient strategies for a stutter jump like the one shown in Fig. 3. Here, the pogo-stick is allowed to leave the ground once more than in the case of a single jump, allowing it to compress the spring farther and jump higher in the final jump. In this jumping strategy, the actions the agent takes early in a jumping episode have a greater effect on the final jump. This is because the more energy it can store in the spring, the higher it will be able to jump.

These two jumping types are chosen based on their relevance in studies using the pogo-stick environment (Vaughan (2013)). Additionally, there are different control

Table 2. Agent Training Schedule

Task	Jump Type	# Agents	Seeds
Jump	Stutter Jump,	10 per	9, 16, 104,
High	One Jump	Jump Type	107, 250,
Conserve	Stutter Jump,	10 per	676, 767,
Power	One Jump	Jump Type	868, 878, 918, 947

Table 3. Pogo-stick Initial State

State Variable	Value
Position of rod, $x$	$-\left[\frac{m_t g}{k}\right]^{1/3}$ mm
Velocity of rod, $\dot{x}$	0.0 m/s
Acceleration of rod, $\ddot{x}$	0.0 m/s <sup>2</sup>
Position of actuator, $x_a$	$\frac{1}{2}(x_a)_{\max}$ mm
Velocity of actuator, $\dot{x}_a$	0 m/s
Acceleration of actuator, $\ddot{x}_a$	0 m/s <sup>2</sup>

strategies that the agent could learn to maximize height. As an example, it might be expected that the agent will, in both jumping instances, learn to initially compress the spring, then launch the pogo-stick. However, the environment is initialized at equilibrium so that the spring is compressed due to the weight of the system. As such, a higher jumping solution will likely take advantage of the initial spring compression. Instead of initially compressing the spring farther, the agent should actuate the system upward until the spring is at its upstretched limit but before the system leaves the ground so that it is not considered airborne and jumping. Once in this position, the agent should proceed accelerating the system downward do to compressing the spring farther than it would have if it started by accelerating downward.

## 6.2 Agent Training

To find control solutions for each task, ten different agents were trained on each task individually. The training schedule can be seen in Table 2. Note that the seeds represent the integer seeds passed to the training algorithm which are used to initialize the neural networks. The same networks are therefore initialized for both high jumping agents and power conserving agents.

During training, the agents were allotted 500,000 steps in their environment, where at each step the agent was presented with the state of the pogo stick which it used to take an action in the environment. The state, and action,  $s_i$ , can be written as:

$$\mathcal{S} = [x_{at}, \dot{x}_{at}, x_t, \dot{x}_t] \quad (7)$$

$$\mathcal{A} = [\ddot{x}_{at}] \quad (8)$$

where,  $x_t$ ,  $\dot{x}_t$ ,  $x_{at}$  and  $\dot{x}_{at}$  are the positions and velocities of the rod and actuator, respectively, and  $\ddot{x}_{at}$  is the acceleration of the actuator.

The environments were initialized using the initial conditions shown in Table 3. The rods starting position was determined by the amount the spring compresses according to the total weight of the system. The actuator initial

position was at the midpoint of its stroke so that it was free to move in either direction.

Episode terminations were defined based on the two different jump types. The first episodic termination stipulation was after the agent completed a single jump. This is defined as the rod's position being greater than zero, then returning to zero. At that point in time, the episode was terminated, and the environment's state was reinitialized. The second type of episode termination was based on the stutter jump motion. This allowed the rod's position to be greater than zero twice, therefore allowing for two jumps. If neither of the two defined terminations occurred, the episode was terminated after 500 time steps, or 5 seconds.

## 7. AGENT EVALUATION

During the evaluation of the 40 different agents that were trained, some of them did not learn good control strategies. These did not learn how to actuate the system in a way which is required to launch the pogo-stick off the ground. This appears to be a neural network initialization issue, as many of the agents which did not get off the ground had the same seed for initializing their networks. Because of this, the agents which did not learn to get off the ground were removed from the evaluation data set. Roughly the same number of agents tasked with power conservation did not succeed at getting off the ground as agent who's task was to jump high.

Figs. 4–5 show time series data of the best performing single jumping agents based on their reward function for both types of agents. These two agents were both trained to actuate the pogo-stick for a single jump. Fig. 4 shows that the agents achieved similar heights, where the high

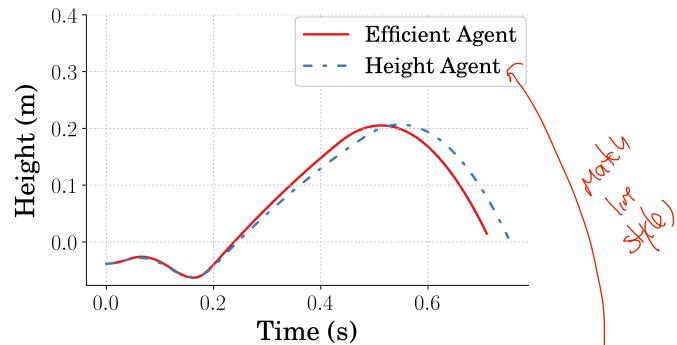


Fig. 4. Example Single Jump Responses

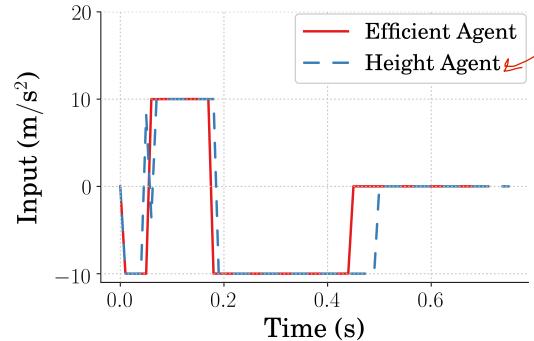


Fig. 5. Example Single Jump Control Inputs

jumping agent jumped only 0.77% higher. Fig. 5 shows that the control input is where there exists a major difference between the agents. The difference in the timing that the two agents learn, lead to the efficient agent being 15.42% more efficient. It has been thoroughly studied that the timing input is important when maximizing jumping height is the goal (Vaughan (2013)). However, it is equally important when seeking to maximize power efficiency.

Figs. 6–7 show time series data of the best performing stutter jumping agents based on their reward function for both types of agents. In the case of the stutter jumping agents, the efficient agent learned both a higher jumping strategy and a more optimal control input. The efficient agent learned a control strategy that jumped 3.78% higher while also being 23.62% more energy efficient. How the efficient agent learned to jump higher can be explained by the fact that higher jumps do increase that agents reward. Reducing power also increases its reward. In this study, for the stutter jumping task, the efficient agent not only learned to jump higher, but also learned to do so while using less power.

Figs. 8–9 show average and best agent performance measures, comparing both agent types that learned strategies for a single jump. When comparing the agents, the average and maximum height reached were relatively similar, with the height agents having jumped 0.61% and 0.76% higher, respectively. This can be explained in that both agent types are seeking to maximize height, but the efficient agents are doing so in a more efficient manner. Looking at Fig. 9 However, it is apparent there is a significant difference in efficiency. Here the efficient agents learned

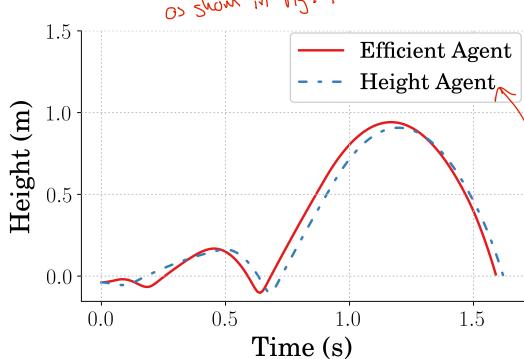


Fig. 6. Example Stutter Jump Responses

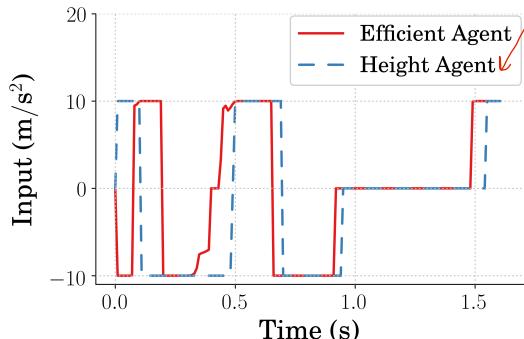


Fig. 7. Example Stutter Jump Control Inputs

Reviewer Q:  
How do you define optimality here?  
Can you prove that optimality?

Remember that optimal ≠ good.  
It has a very specific meaning in  
o controls & RL context.

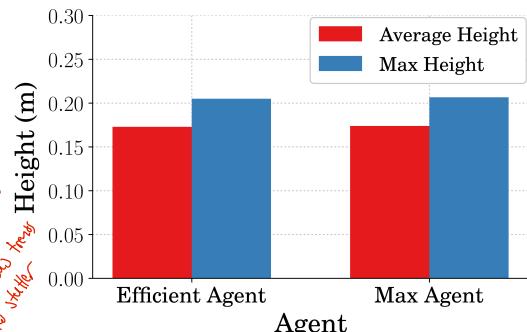


Fig. 8. Jump Heights for Stutter Jumps

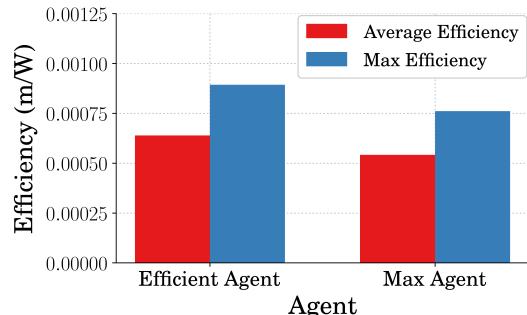


Fig. 9. Efficiency for Stutter Jumps

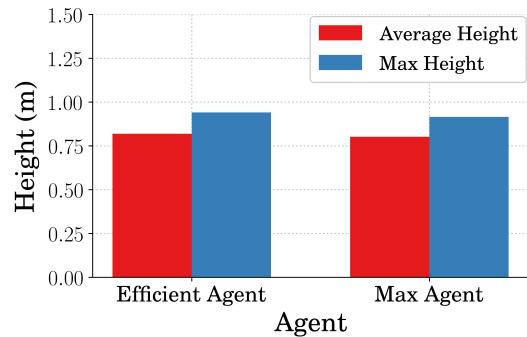


Fig. 10. Jump Heights for Stutter Jumps



Fig. 11. Efficiency for Stutter Jumps

(average and maximum efficient strategies that are 17.91% and 17.33% more efficient, respectively.)

Similar to the case of the single jumping agents, data was collected and averaged for the stutter jumping agents and is shown in Figs. 10–11. Again, similar to the single jumping agent, the stutter jumping agents average and

outward walking.  
They are the same strategy, just different measures.

max heights reached are close, only this time the efficient agents managed higher jumping performance by 2.22% and 2.75%, respectively. Looking at the efficiency data for the stutter jumping agents, the efficient agents not only learned higher jumping strategies, but also more efficient ones. The average and maximum efficiencies difference is 10.56% and 13.84%, respectively. The implications from this data match those of the times series data presented in Figs. 6–7. That is, in the case of learning to stutter jump, the efficient agents not only outperform the high jumping agents in terms of efficiency but also in terms of height reached.

## 8. CONCLUSION

Two different types of agents were trained on the pogo-stick environment meant to represent a flexible-legged jumping system. The first type of agents were trained and deployed to jump high, and the second type were trained and deployed to jump high while minimizing power. Both types of agents were trained to jump in two different ways; one to jump once and one to stutter jump. The results presented show that in the case of the single jumping agents performance, the efficient agents learn strategies that jump slightly lower but use significantly less power. In the case of the stutter jumping agents performance, the efficient agents learn both higher jumping and more efficient strategies. The general implications of the results presented are that RL control strategies can be trained to learn more efficient strategies with a properly defined reward function.

## ACKNOWLEDGEMENTS

Thanks to the Louisiana Crawfish Board for their support in the CRAWLAB. It is through their support that research like this is possible.

## REFERENCES

- Ahmadi, M., Buehler, M., 1997. Stable control of a simulated one-legged running robot with hip and leg compliance. *IEEE Transactions on Robotics and Automation* 13 (1), 96–104.
- Blackman, D. J., Nicholson, J. V., Pusey, J. L., Austin, M. P., Young, C., Brown, J. M., Clark, J. E., 2018. Leg design for running and jumping dynamics. *2017 IEEE International Conference on Robotics and Biomimetics, ROBIO 2017* 2018-Janua, 2617–2623.
- Blickhan, R., Full, R. J., 1993. Similarity in multilegged locomotion: Bouncing like a monopode. *Journal of Comparative Physiology A* 173 (5), 509–517.
- Cui, L., Chen, W., Wang, H., Wang, J., 2019. Control of Flexible Manipulator Based on Reinforcement Learning. *Proceedings 2018 Chinese Automation Congress, CAC 2018*, 2744–2749.
- Dwiel, Z., Candadai, M., Phielipp, M., 2019. On Training Flexible Robots using Deep Reinforcement Learning. *IEEE International Conference on Intelligent Robots and Systems*, 4666–4671.
- Folkertsma, G. A., Kim, S., Stramigioli, S., 2012. Parallel stiffness in a bounding quadruped with flexible spine. *IEEE International Conference on Intelligent Robots and Systems*, 2210–2215.
- Fujimoto, S., Van Hoof, H., Meger, D., 2018. Addressing Function Approximation Error in Actor-Critic Methods. *35th International Conference on Machine Learning, ICML 2018* 4, 2587–2601.
- Galloway, K. C., Clark, J. E., Yim, M., Koditschek, D. E., 2011. Experimental investigations into the role of passive variable compliant legs for dynamic robotic locomotion. *Proceedings - IEEE International Conference on Robotics and Automation*, 1243–1249.
- Harper, M. Y., Nicholson, J. V., Collins, E. G., Pusey, J., Clark, J. E., 2019. Energy efficient navigation for running legged robots. *Proceedings - IEEE International Conference on Robotics and Automation 2019-May*, 6770–6776.
- He, W., Gao, H., Zhou, C., Yang, C., Li, Z., 2020. Reinforcement Learning Control of a Flexible Two-Link Manipulator: An Experimental Investigation. *IEEE Transactions on Systems, Man, and Cybernetics: Systems*, 1–11.
- Hurst, J., 2008. The Role and Implementation of Compliance in Legged Locomotion. *The International Journal of Robotics Research* 25 (4), 110.
- Iida, F., Gómez, G., Pfeifer, R., 2005. Exploiting body dynamics for controlling a running quadruped robot. *2005 International Conference on Advanced Robotics, ICAR '05, Proceedings 2005*, 229–235.
- Park, H. W., Wensing, P. M., Kim, S., 2017. High-speed bounding with the MIT Cheetah 2: Control design and experiments. *International Journal of Robotics Research* 36 (2), 167–192.
- Peng, X. B., Berseth, G., Panne, M. V. D., 2016. Terrain-Adaptive Locomotion Skills Using Deep Reinforcement Learning 35 (4), 1–12.
- Pratt, G. A., Williamson, M. M., 1995. Series elastic actuators. *IEEE International Conference on Intelligent Robots and Systems* 1, 399–406.
- Reda, D., Tao, T., van de Panne, M., 2020. Learning to Locomote: Understanding How Environment Design Matters for Deep Reinforcement Learning.
- Seok, S., Wang, A., Chuah, M. Y., Hyun, D. J., Lee, J., Otten, D. M., Lang, J. H., Kim, S., 2015. Design principles for energy-efficient legged locomotion and implementation on the MIT Cheetah robot. *IEEE/ASME Transactions on Mechatronics* 20 (3), 1117–1129.
- Seok, S., Wang, A., Chuah, M. Y., Otten, D., Lang, J., Kim, S., 2013. Design principles for highly efficient quadrupeds and implementation on the MIT Cheetah robot. In: *Proceedings - IEEE International Conference on Robotics and Automation*. IEEE, pp. 3307–3312.
- Sugiyama, Y., Hirai, S., 2004. Crawling and jumping of deformable soft robot. *2004 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS) 4 (c)*, 3276–3281.
- Thuruthel, T. G., Falotico, E., Renda, F., Laschi, C., 2019. Model Based Reinforcement Learning for Closed Loop Dynamic Control of Soft Robotic Manipulators. *IEEE Transactions on Robotics* 35, 124–134.
- Vaughan, J., 2013. Jumping Commands For Flexible-Legged Robots.
- Yang, Y., Caluwaerts, K., Iscen, A., Zhang, T., Tan, J., Sindhwani, V., 2019. Data efficient reinforcement learning for legged robots. *arXiv (CoRL)*, 1–10.